Computer Vision Layer for Robust Model

Group "Jean Ponce" :

Gabriella FERNANDES MACEDO, Ruben IFRAH,
Clément ROUVROY
Dauphine-PSL, École Polytechnique, ENS-PSL

Intuition : Simplification as Defense

- Attack : Adds high-frequency noise (textures, details).
- Defense : Remove details, keep structure.

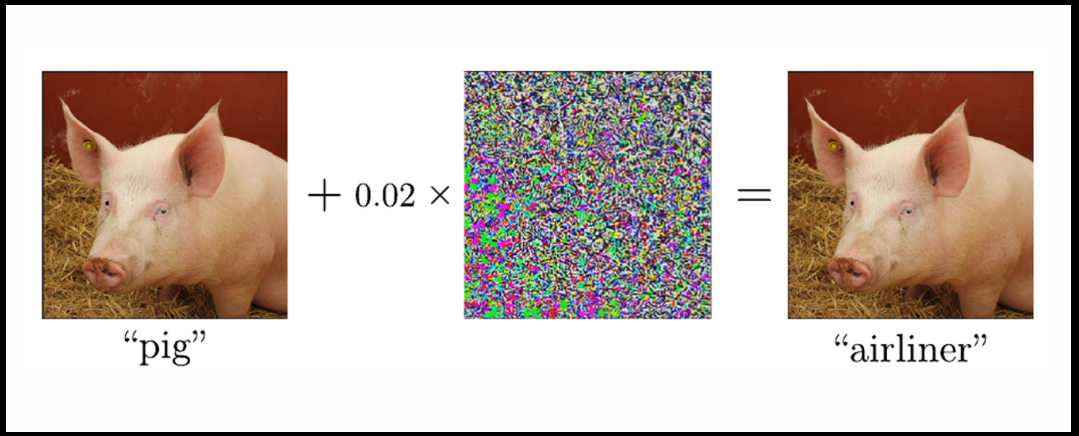

Example : Pig classified as Airliner

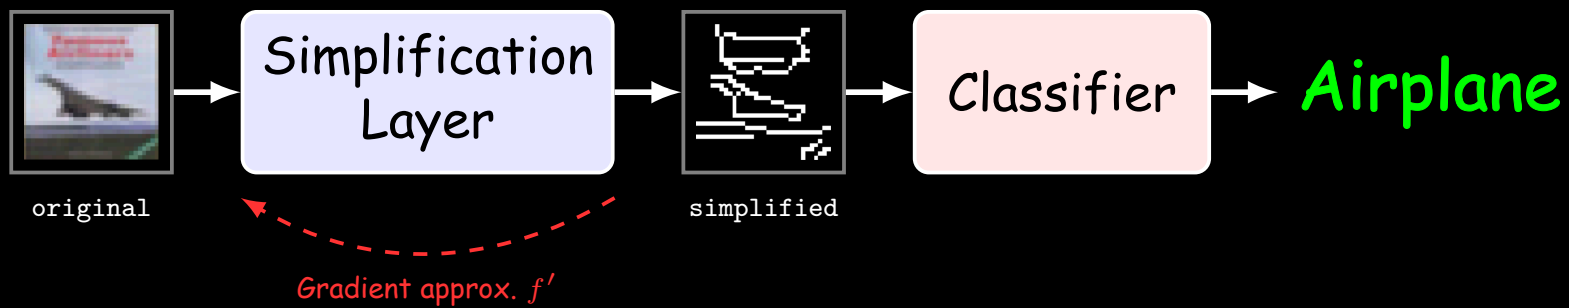

Canny Edge : Mechanism

1. Gaussian Blur (Denoise)
   - Removes high-freq noise.
   - Prevents false detections.

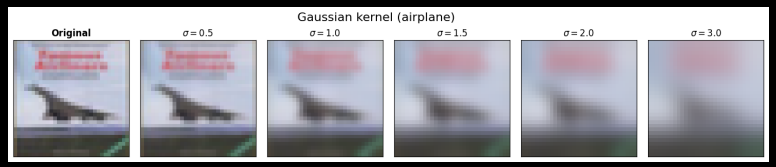

2. Gradient Calculation
   - Sobel operators (G_x, G_y).
   - Magnitude & Direction.

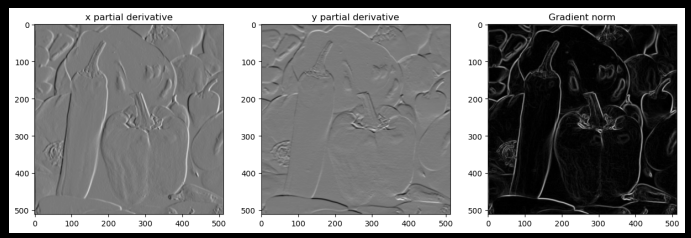

3. Thresholding : Hysteresis to connect weak edges to strong ones.

Canny Edge : Implementation

Implementation Details :

- Library : `kornia.filters.canny`
- Hyperparameter Tuning :
  - Calibrated on CIFAR-10.
  - Goal : "Is object visible?" vs "Too much noise?".

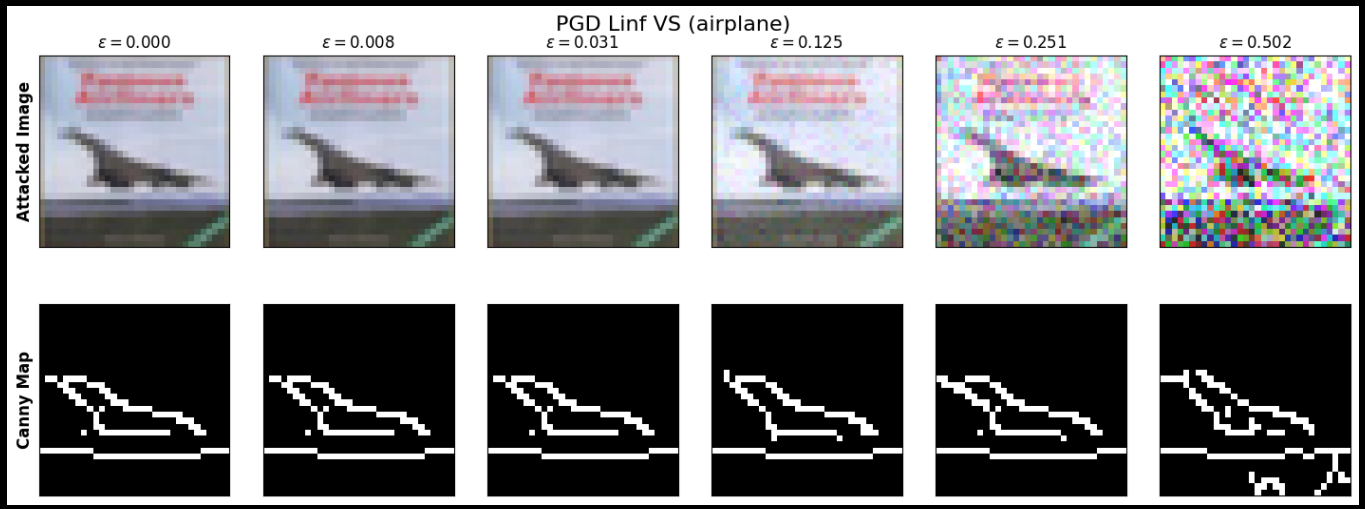

Attacked Image vs. Canny Output

Other Non-Differentiable Methods

Explored Methods :

- Quantization : Discretizing pixel values.
- Mean-Shift : Clustering-based smoothing.
- Median Filtering : Noise removal.
- Combinations : Stacking filters.

Why they failed (vs. Canny) :

- Not destructive enough : They preserve too much structure/linearity.
- Ineffective against L1 attacks : Sparse but high-magnitude noise survives these filters.
- *Canny Edge is a radical simplification that neutralizes texture attacks.*

Data Augmentation

Impact on generalization :

- Without it :
  Natural Acc $\approx$ 68%.
- With it :
  Natural Acc $\rightarrow$ 100%.

Techniques :

- Random Crop : Pad 4px + Crop 32x32.
- Random Horizontal Flip : p=0.5.

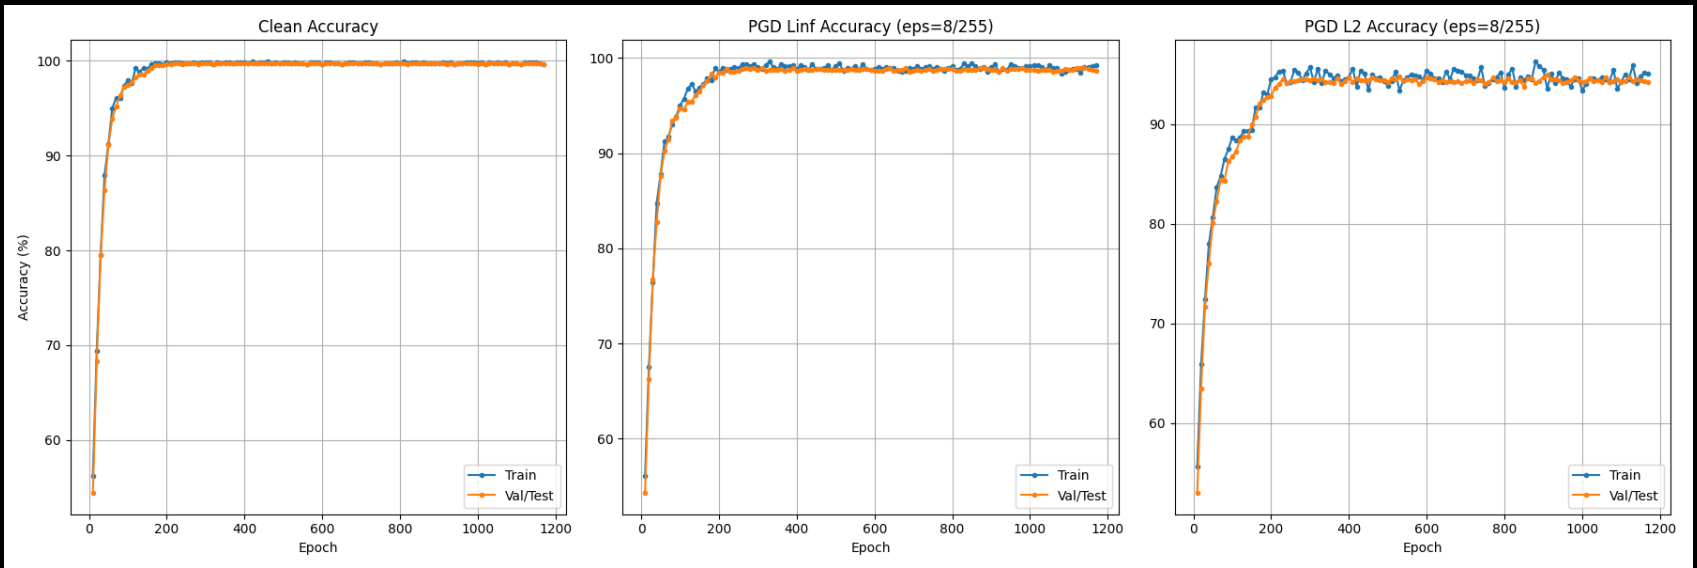

The "Trick" : Gradient Obfuscation

To reach 197.29/200 robustness score :

Alternating Gradient in Backward Pass :

- Instead of Identity, we return $+1$ or -1.
- Breaks the PGD optimizer (Gradient Obfuscation).

```
def backward(ctx, grad_output):
    if call_count % 2 == 0: return 1.0
    else: return -1.0
```

Note : Effective against PGD, but not true robustness (adaptive attacks would work).

Results

Using Canny Edge Filter + Data Augmentation :

- Nat acc : 100%
- L inf acc : 60.53% L2 acc : 89.59%

Using "the trick" :

VICTORY ROYAL

- Natural Accuracy : 100%
- Robustness Score : 197.29 / 200

Thank you for your attention

BONUS

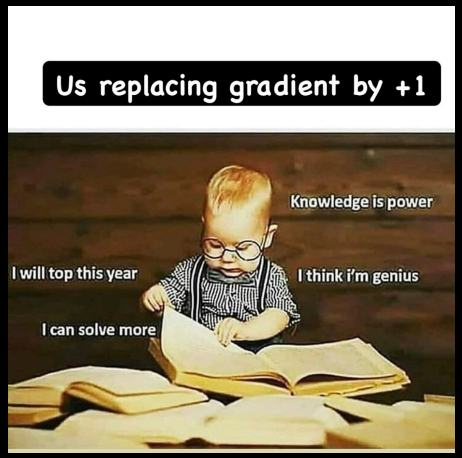
Us replacing gradient by +1
Knowledge is power
I will top this year
I think i'm genius
I can solve more


JEAN PONCE
CLASSIC BEAMER PRESENTATION
OR
draw 25


Our model
"No idea"
Our model
"An airplane"

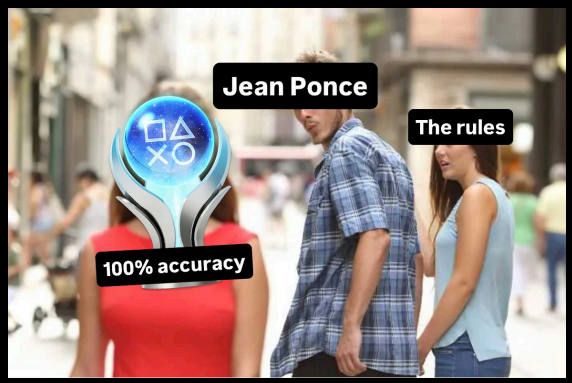
Jean Ponce
The rules
100% accuracy

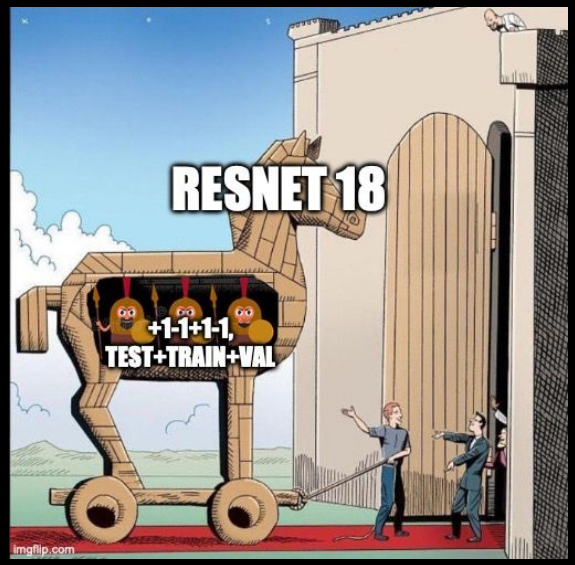
RESNET 18
+1-1+1-1,
TEST+TRAIN+VAL
imgflip.com


VERRINE
NOTRE CODE